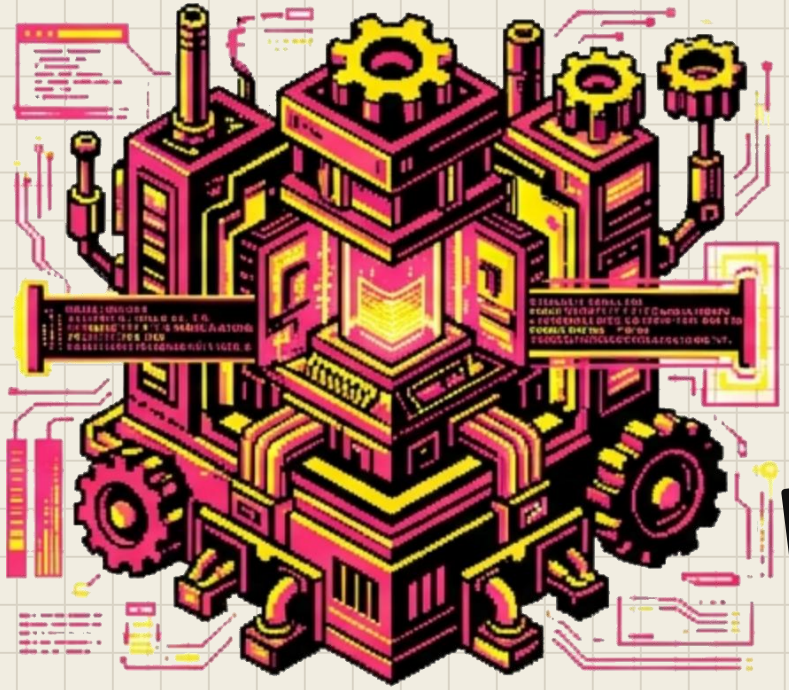
LOCH RAVEN CAPITAL

V 2.0 // AN AI FIRM WITH INSTITUTIONAL MEMORY

An AI investment firm that remembers what it got right, and what it got wrong.



Claude research floor
Hybrid RAG memory
RAFT fine tuning data
Hard risk limits

HACKUMBC 2026
T. ROWE PRICE TRACK

Patrick Kamdem
Morgan State University

SEPTEMBER 10, 2026

T. Rowe Price brings Claude into its investment research.

“...where the bar for accuracy and judgment is highest.”

Ramon Richards, Chief Technology Officer, T. Rowe Price, on bringing Claude into the firm’s research. Press release, September 10, 2026.



LOCH RAVEN RUNS ON CLAUDE

Five research desks, a debate room and a portfolio manager. Every one of them is Claude.

When an AI helps make an investment decision, how do you know it was right?

CONFIDENT IS NOT THE SAME AS RIGHT

01

GENERIC MEMORY

It answers from what the model read online, not what your firm already learned.

02

NO SOURCES

It cannot show which evidence drove the call, so nobody can audit it.

03

NEVER GRADED

Nobody checks whether its past calls were right, so it never gets better.

41%

of my own AI committee's 264 buy and sell calls were right, measured against the S&P 500 over the next 20 trading days.

1,794 decisions graded / June to September 2026

THE AI TEAM. AN AGENTIC FLOW.

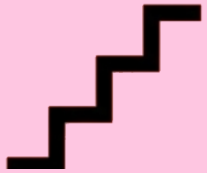
BUILT BEFORE THIS WEEKEND // JUNE TO SEPTEMBER 2026 // RUNS UNATTENDED



QUANT SCREEN
Every listed US stock, every morning



5 RESEARCH DESKS
Fundamental, technical, news, macro, quant



DEBATE ROOM
Bull, bear and base case attack each other



PM JUDGE
Binding rating, conviction and size



RISK DESK
Hard limits and a kill switch



EXECUTION
Four strategy sleeves, $100k paper book

[1,104 regression checks pass / 5 audits / Claude on every desk]

EVERY DESK IS CLAUDE.

FOUR SLEEVES. $100K. ONE AI TEAM.

PAPER BOOK SINCE JUNE 9, 2026 // $100,309 AT THE SEPTEMBER 25 CLOSE



ALPHA STOCKS

Committee sized single stock book. Buys at conviction 60 plus, 18% cap per name.

START $25,000

NOW $22,651

minus 9.4% since June



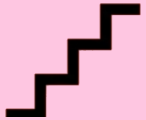
CORE BLEND

55% ETF core earned by momentum, 45% conviction satellites.

START $25,000

NOW $26,098

+4.4% since June



OPTIONS DESK

Long calls and puts only. Defined risk, AI designed exit ladders.

START $25,000

NOW $25,450

+1.8% since June



ETF ROTATION

Top 3 of 17 ETFs by momentum each month. T bills when nothing beats cash.

START $25,000

NOW $26,109

+4.4% since June

HOW THE AI TEAM ALLOCATES

PM SIZES
every position by conviction

→

RISK DESK CAPS
18% a name, 35% a sector

→

WEEKLY REBALANCE
cash moves when a sleeve leaves 12 to 40% of the firm

→

LEARNING REVIEW
retunes caps inside hard bounds

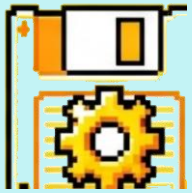
Rule based sleeves tracked the S&P 500, up 4.6%. The sleeves that lean on AI judgment lagged. So I built memory.

THIS WEEKEND. INSTITUTIONAL MEMORY.

BUILT AT HACKUMBC 2026

THE PLAYBOOK

how we analyze

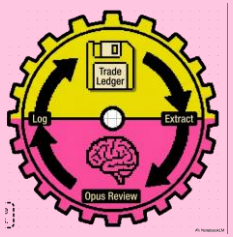


Nine chapters the AI can search. DCF, multiples, thesis work, statements, modeling and risk policy, plus 12 lessons the firm learned from its own trades.

Written from my Wall Street Prep equity research certification and a DCF lesson series a T. Rowe Price investment liaison taught at Morgan State.

THE HISTORY

what we decided



1,794 committee decisions, each graded against the S&P 500 twenty trading days later.

Point in time. A decision only sees memos written before it and outcomes that were already known.

RETRIEVE
before every verdict

→

CITE
every claim as K# or P#

→

CHECK
citations and numbers

→

DECIDE
the same risk desk trades it

LIVE. SAME COCKPIT. MEMORY INSIDE.



1 LAST NIGHT'S RUN
five names, memory on. The ETF QQQ held at 45, no chasing

2 CITED MEMORY
P2 our June QQQ buy, graded wrong. K2 firm lesson 6

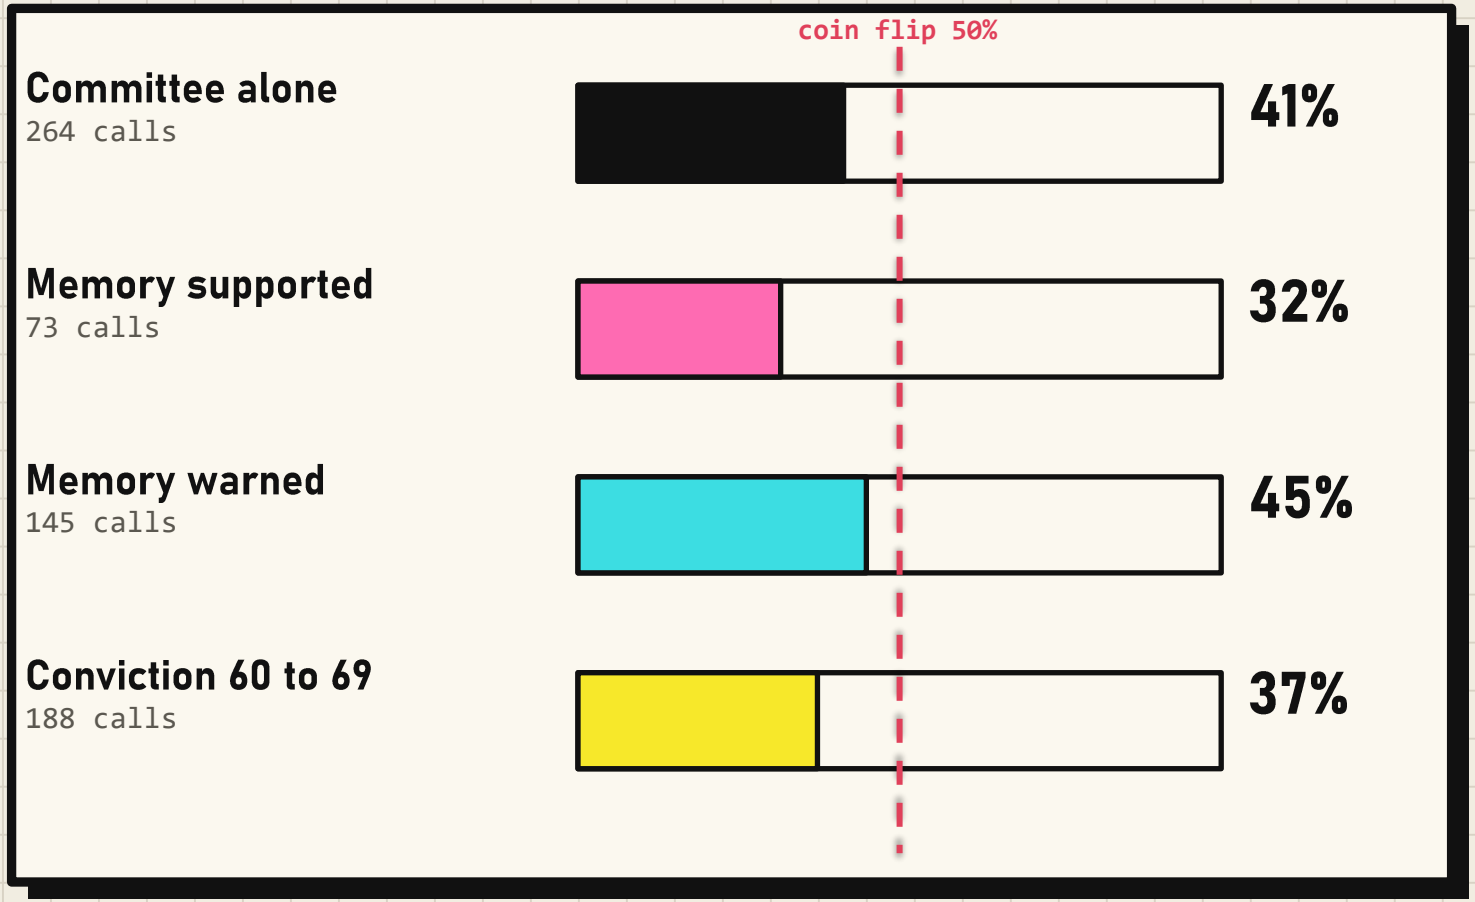
3 CHECKED
citations real, all 14 figures traced, logged to the audit chain

4 SAME FIRM
four sleeves, $100k paper, the same risk desk

[the original cockpit / paper copy of the live books / committee run of September 27 with memory on]

THE NUMBER MOST AI PITCHES HIDE

Walk forward test over every graded buy and sell call. Each call only sees what was knowable on its own date.



SO MEMORY INFORMS. IT NEVER OVERRULES.

Momentum flipped in July, so setups that worked in June stopped working.

IT FOUND A REAL FLAW

71% of calls sat at 60 to 69 conviction and were right 37% of the time. The firm's own review flagged it on September 18. Now it is cited at every decision.

[6 precedents per call, as live / at depths 5 to 12 supported calls were right 32 to 41%, warned calls 40 to 45% / 218 calls, not significant]

RAG GIVES KNOWLEDGE. FINE TUNING SHAPES JUDGMENT.

RAFT FORMAT

Every example carries the playbook passages and the precedents it could have retrieved that day, so the model learns to use the firm's memory.

TIME SPLIT

Trained only on calls graded before August 10, tested on calls made after. 0 leaked rows.

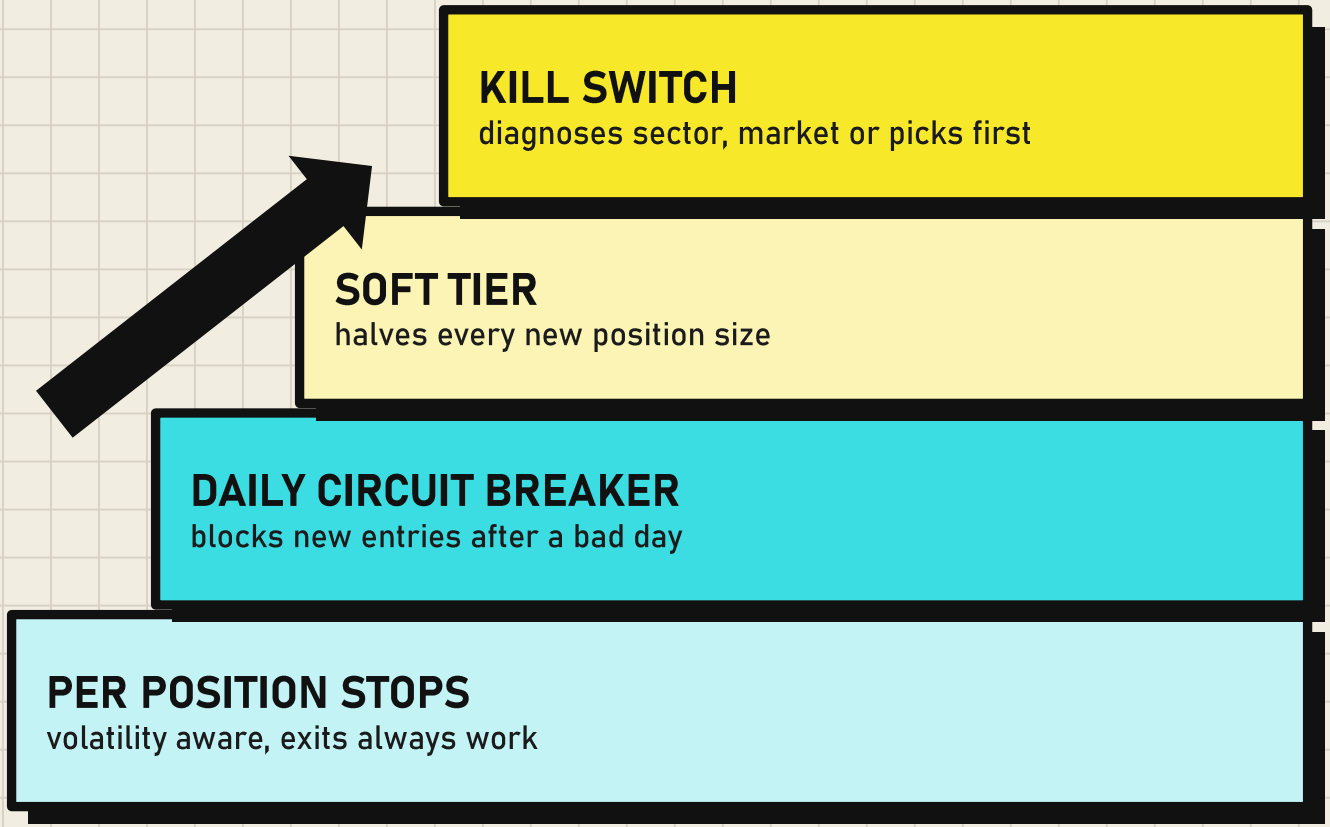
LABELED BY THE MARKET

Desirable if the call beat the S&P 500 in its direction, undesirable if it did not.

DATASET	TRAIN	TEST
SFT, proven right only	15	17
KTO, labeled by outcome	53	48
Reward labels	1,410	
Leaked rows	0	

Four months gives a starter set. Ninety years gives a real one.

RESPONSIBLE BY DESIGN



EVERY ANSWER IS CHECKED

- Only real citations
- Every number checked against evidence
- No look ahead, ever
- Firm limits the model cannot move
- Paper money until a human switches
- Tamper evident audit chain

The AI reasons inside the system. It cannot reason its way around it.

SAME AGENTIC FLOW. YOUR QUALITY DATA.

MY ARCHIVE

4 months

1,794 decisions / 255 tickers



T. ROWE PRICE

nearly 90 years

of research / $1.87 trillion in client assets

T. Rowe Price press release, September 10, 2026. Client assets as of July 31, 2026.

RAG gives AI access to institutional memory.

AGENTS give that knowledge a process.

OUTCOMES make the process learn.

GUARDRAILS keep it responsible.

LOCH RAVEN CAPITAL

Patrick Kamdem
Morgan State University / Finance and Information Systems
Certified equity researcher / System specialist, Morgan State Food Resource Center

THANK YOU.
QUESTIONS?

github.com/PadiusD1/loch-raven-capital

[Claude / hybrid RAG over 1,794 graded decisions and a 9 chapter playbook / RAFT data / 1,221 automated checks]